

ASHUTOSH TIWARI

669-292-7534 | findashutoshtiwari@gmail.com | [linkedin.com/in/ashutosh--tiwari](https://www.linkedin.com/in/ashutosh--tiwari) | thunderrock.github.io | github.com/thunderrock

WORK EXPERIENCE

Senior Machine Learning Engineer (ML Platform & Frameworks)

Jul. 2024 – Present

ADOBE FIREFLY

GENERATIVE AI, COMPUTER VISION, LARGE LANGUAGE MODELS

- Lead fault-tolerant distributed inference and training infrastructure for the org's ML platform: GPU pipelines on Kubernetes engineered to survive node loss and partial failure, from 40-node daily EMR ingestion through production model inference.
- **Distributed Inference Frameworks:** Designed and led engineering for the org's online and offline inference systems powering experimentation and production at scale.
 - Built online inference on vLLM + Ray for rapid experimentation with LLMs such as Falcon-40B and Llama 2 70B.
 - Built offline inference on Ray Data + PyTorch Lightning, achieving 10–50× faster image loading and 3–8K images/sec on 32 GPUs.
 - Modeled pipelines as horizontally-scaling Source → Transform → Sink plugins with no pod-to-pod communication, KEDA-autoscaled on SQS queue depth.
- **Foundation Model Training Framework:** Leading development of a PyTorch-based training framework adopted across the org for training generative models on billions of images and videos.
 - Designed a plugin-based strategy pattern over FSDP/FSDP2 to make distributed training strategies pluggable.
 - Integrated Flash Attention 2/3, context parallelism, and distributed attention for DiT text-to-image and text-to-video training.
- **Build & Deployment System:** Built a CLI-driven build system (BuildKit images on KEDA-scaled Kubernetes jobs, pushed to ECR), fronted by a Rust control plane and an MCP interface for AI-agent-driven builds, deployments, and inference.
- **Enrichment Service:** Orchestration and abstraction layer over the inference framework and other ML Platform offerings; runs fire-and-forget pipelines on datasets of billions of images and videos.
- **Training-Data Governance:** Solely authored the org's fault-tolerant training-data lineage subsystem in Rust (27,135 lines): idempotent claim/receipt registration, immutable S3 evidence layouts, and manifest-after-data publication, proven by a 9.7k-line integration test.
- **Data Quality Framework:** Wrote the org's first data quality framework on Cerberus, used across enrichment pipelines to validate the correctness of feature-generation outputs.

Software Dev Engineer II (ML Platform)

Jan. 2019 – Jul. 2021

SWIGGY

DATA SCIENCE PLATFORM, TIME SERIES FORECASTING

- **Online Inference Platform:** Core contributor to the online inference platform serving 18M+ monthly transacting users across food delivery and quick commerce, with models autoscaled by real-time load metrics.
 - Engineered a dynamic request-routing layer that evaluated incoming API calls against geo-tags, traffic profiles, and service metadata, then dispatched them to model clusters partitioned by region and load domain.
 - Implemented routing/dispatch on an Akka-style actor framework for resilient asynchronous invocation, fault isolation across clusters, and backpressure-aware scheduling at sustained low P99 under high QPS.
- **Feature Store:** Built and maintained a Feature Store and pipeline serving on-demand features to deployed ML models at production scale (4Bn rows, 10K QPS), with multichannel ingestion via Spark, Flink, and user files.
- **Forecasting & Correlation Platform:** Founding member of the platform used by many teams to forecast time series; forecasts powered critical scaling decisions across the org in real time.
- **DAQ:** Led DAQ, a tool used to scrape APIs at scale, collecting 15M rows daily for analysis and model training.

Software Development Engineer (Search Relevance)

Sep. 2017 – Jan. 2019

FLIPKART (a Walmart company)

SEARCH RELEVANCE, QUERY INTENT, NLP

- **Search Intent Models:** Built and improved CRF/Neural Network-based intent models powering user search and discovery for millions every day; identified error classes, designed solutions, and shipped fixes.
- **Tail Query Classifier:** Implemented a FastText-based query store classifier that predicts the category of a tail query.
- **ML Workflow Automation:** Built the first workflow at Flipkart to automate training and auto-deployment of search models—written first in Luigi and later migrated to Airflow.
 - Wrote a generic Airflow framework that creates DAGs at runtime for different ML models and orchestrates the train→deploy flow, including data and model validations.
- **Large-Scale Data Pipelines:** Implemented Cascading/HDFS pipelines (4Bn+ datapoints) to extract user-event data and transform it for model training.

Software Development Engineer

Sep. 2016 – Sep. 2017

GROUPON

BACKEND ENGINEERING

- Customer segmentation and deal-recommendation ML (NSVD unsupervised collaborative filtering on Neo4j); Cyclops customer-service platform live in every Groupon country.

Software Engineer

Sep. 2015 – Aug. 2016

NETSPEED SYSTEMS (Acquired by Intel)

GRAPH ALGORITHMS, NETWORK ON CHIP

- C++ graph algorithms for Network-on-Chip interconnects: Virtual Channel Arbitration, Polarity-based Arbitration, Multi-Cast Filtering, Structural Latency Breakdown.

PUBLICATIONS

Accepted at [NetSci 2023](#) (Poster Presentation) and [IC2S2 2023](#) (Parallel Talk) as **first author**

- [1] **Ashutosh Tiwari**, Prof. Sadamori Kojaku, Prof. Yong-Yeol Ahn, "Biased Contrastive Learning debiases Graph Neural Networks," *International Conference on Network Science (NetSci)*, 2023. [Poster](#)

In this work, we propose a non-parametric contrastive learning framework to learn debiased graph embeddings with respect to sensitive node attributes and structural homophily. Through empirical evaluations on different datasets, we demonstrate that our method offers a better approach to debiasing compared to existing approaches and thus results in more organic recommendations across different GNN architectures.

EDUCATION

Indiana University, Bloomington

Master of Science in Computational Data Science 3.87/4.0

Aug. 2021 – May 2023

National Institute of Technology

Bachelor of Science in Computer Science 8.32/10.0

Aug. 2011 – May 2015

RESEARCH EXPERIENCE / INDEPENDENT STUDY

- Implementing novel training frameworks that result in unbiased (or fair) recommendation systems from "biased" datasets at [IUNI](#) with Prof. YY Ahn and Prof. S Kojaku as part of my Independent Study. Summer 2022 - Summer 2023
- Working as a paid RA for, "**User Intent as a Network**" with Prof. YY Ahn, P Kantak and FB Yara. Project is funded by Kelly Business School. Goal is to be able to quantify and thus act upon that "intent" network. Summer 2022 - Fall 2022
- Was part of [NLP Lab@IUB](#) with Prof. D Cavar. Contributed to design of [TieML](#) and Events' Timeline modelling using different fine-tuned Large Language Models (LLMs). Fall 2021

TEACHING EXPERIENCE

- Teaching Assistant for [Network Science](#) (INFO-I 606) with Prof. YY Ahn Spring 2023
- Teaching Assistant for [Machine Learning](#) (CSCI-B 555) with Prof. R Khardon Fall 2022
- Teaching Assistant for [Network Science](#) (INFO-I 606) with Prof. YY Ahn Spring 2022

SELECTED PROJECTS

- BiasNet** | *Deep Reinforcement Learning, PyTorch, Actor Critic Algorithm, On-policy Model Free* [Code/Report](#)
– Learning to fight in [Street Fighter II](#) with induced relational bias from differential game scenes.
- DeepFoodie** | *Python, Tensorflow, Self Supervised Deep Clustering, Deep Learning, Transfer Learning* [Code/Report](#)
– Clustering dishes on basis of their ingredient embeddings. These ingredient embeddings are generated by a NN.
- BlindNet** | *Python, Pytorch, Deep Learning, Transfer Learning* [Code/Report](#)
– Image to vector generation on Coco dataset.
- Continuous Dominant Set Repair** | *C++, Graph Algorithms, Guha and Khuller's Algo., Greedy* [Code](#)
– Repairs a broken link in Continuous Dominant Set in $O(\Delta^2)$, where Δ being the avg cardinality of connected graph.
- Humana Mays Healthcare Analytics Case Competition** | *Boosted Trees, Feature Engineering* [Code](#)
– 11th [rank](#) on leaderboard. Hosted by TAMU and Humana Mays, 2021.
- AnalyticsVidhya Solutions** | *Python, Tensorflow, Torch, Time Series, Feature Engineering, Catboost* [Code](#)
- Investigating Bias Manifolds** | *Bias Manifolds, Bias Progression, Measuring Bias, Python, Word2vec* [Code/Report](#)

TECHNICAL SKILLS

Search Relevance, Recommendations, Graph Neural Networks, Fairness Aware Modeling, Computational ML, NLP, Computer Vision, DL, Spark, Python, Scala, C++, Contrastive Learning, Large Language Models

Frameworks/Libraries/Tools: Ray, Pytorch Geometric, Pytorch Lightning, Pytorch, Tensorflow, Sklearn, Numpy, Pandas, Matplotlib, HDFS, Spark, Kafka, Flink, AWS Sagemaker, DynamoDB, Faiss, vLLM, Django